

# Lab 2: Cats vs Dogs

**Deadline:** Feb 01, 5:00pm

**Late Penalty:** There is a penalty-free grace period of one hour past the deadline. Any work that is submitted between 1 hour and 24 hours past the deadline will receive a 20% grade deduction. No other late work is accepted. Quercus submission time will be used, not your local computer time. You can submit your labs as many times as you want before the deadline, so please submit often and early.

**Marking TA:** Tinglin (Francis) Duan

This lab is partially based on an assignment developed by Prof. Jonathan Rose and Harris Chan.

In this lab, you will train a convolutional neural network to classify an image into one of two classes: "cat" or "dog". The code for the neural networks you train will be written for you, and you are not (yet!) expected to understand all provided code. However, by the end of the lab, you should be able to:

1. Understand at a high level the training loop for a machine learning model.
2. Understand the distinction between training, validation, and test data.
3. The concepts of overfitting and underfitting.
4. Investigate how different hyperparameters, such as learning rate and batch size, affect the success of training.
5. Compare an ANN (aka Multi-Layer Perceptron) with a CNN.

## What to submit

Submit a PDF file containing all your code, outputs, and write-up from parts 1-5. You can produce a PDF of your Google Colab file by going to **File > Print** and then save as PDF. The Colab instructions has more information.

**Do not submit any other files produced by your code.**

Include a link to your colab file in your submission.

Please use Google Colab to complete this assignment. If you want to use Jupyter Notebook, please complete the assignment and upload your Jupyter Notebook file to Google Colab for submission.

With Colab, you can export a PDF file using the menu option `File -> Print` and save as PDF file.

**Adjust the scaling to ensure that the text is not cutoff at the margins.**

## Colab Link

Include a link to your colab file here

Colab Link: <https://colab.research.google.com/drive/1iDiNPHyq0gmdb2wzeaa1S73US29Ky7QI?usp=sharing> (<https://colab.research.google.com/drive/1iDiNPHyq0gmdb2wzeaa1S73US29Ky7QI?usp=sharing>)

In [ ]:

```
import numpy as np
import time
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
from torch.utils.data.sampler import SubsetRandomSampler
import torchvision.transforms as transforms
```

## Part 0. Helper Functions

We will be making use of the following helper functions. You will be asked to look at and possibly modify some of these, but you are not expected to understand all of them.

You should look at the function names and read the docstrings. If you are curious, come back and explore the code *after* making some progress on the lab.

In [ ]:

```
#####
# Data Loading

def get_relevant_indices(dataset, classes, target_classes):
    """ Return the indices for datapoints in the dataset that belongs to the
    desired target classes, a subset of all possible classes.

    Args:
        dataset: Dataset object
        classes: A list of strings denoting the name of each class
        target_classes: A list of strings denoting the name of desired classes
                       Should be a subset of the 'classes'

    Returns:
        indices: list of indices that have labels corresponding to one of the
                 target classes

    """
    indices = []
    for i in range(len(dataset)):
        # Check if the label is in the target classes
        label_index = dataset[i][1] # ex: 3
        label_class = classes[label_index] # ex: 'cat'
        if label_class in target_classes:
            indices.append(i)
    return indices

def get_data_loader(target_classes, batch_size):
    """ Loads images of cats and dogs, splits the data into training, validation
    and testing datasets. Returns data loaders for the three preprocessed data-
    sets.

    Args:
        target_classes: A list of strings denoting the name of the desired
                       classes. Should be a subset of the argument 'classes'
        batch_size: A int representing the number of samples per batch

    Returns:
        train_loader: iterable training dataset organized according to batch size
        val_loader: iterable validation dataset organized according to batch size
        test_loader: iterable testing dataset organized according to batch size
        classes: A list of strings denoting the name of each class

    """

    classes = ('plane', 'car', 'bird', 'cat',
               'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
    #####
    # The output of torchvision datasets are PILImage images of range [0, 1].
    # We transform them to Tensors of normalized range [-1, 1].
    transform = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
    # Load CIFAR10 training data
    trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                             download=True, transform=transform)
    # Get the list of indices to sample from
    relevant_indices = get_relevant_indices(trainset, classes, target_classes)

    # Split into train and validation
```

```

np.random.seed(1000) # Fixed numpy random seed for reproducible shuffling
np.random.shuffle(relevant_indices)
split = int(len(relevant_indices) * 0.8) #split at 80%

# split into training and validation indices
relevant_train_indices, relevant_val_indices = relevant_indices[:split], relevant_indices[split:]
train_sampler = SubsetRandomSampler(relevant_train_indices)
train_loader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                            num_workers=1, sampler=train_sampler)

val_sampler = SubsetRandomSampler(relevant_val_indices)
val_loader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                          num_workers=1, sampler=val_sampler)

# Load CIFAR10 testing data
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform)

# Get the list of indices to sample from
relevant_test_indices = get_relevant_indices(testset, classes, target_classes)

test_sampler = SubsetRandomSampler(relevant_test_indices)
test_loader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          num_workers=1, sampler=test_sampler)

return train_loader, val_loader, test_loader, classes

#####
# Training
def get_model_name(name, batch_size, learning_rate, epoch):
    """ Generate a name for the model consisting of all the hyperparameter values

    Args:
        config: Configuration object containing the hyperparameters
    Returns:
        path: A string with the hyperparameter name and value concatenated
    """
    path = "model_{0}_bs{1}_lr{2}_epoch{3}".format(name,
                                                  batch_size,
                                                  learning_rate,
                                                  epoch)

    return path

def normalize_label(labels):
    """
    Given a tensor containing 2 possible values, normalize this to 0/1

    Args:
        labels: a 1D tensor containing two possible scalar values
    Returns:
        A tensor normalize to 0/1 value
    """
    max_val = torch.max(labels)
    min_val = torch.min(labels)
    norm_labels = (labels - min_val)/(max_val - min_val)
    return norm_labels

def evaluate(net, loader, criterion):
    """ Evaluate the network on the validation set.

```

**Args:**

*net: PyTorch neural network object*  
*loader: PyTorch data loader for the validation set*  
*criterion: The loss function*

**Returns:**

*err: A scalar for the avg classification error over the validation set*  
*loss: A scalar for the average loss function over the validation set*

```
"""
```

```
total_loss = 0.0
total_err = 0.0
total_epoch = 0
for i, data in enumerate(loader, 0):
    inputs, labels = data
    labels = normalize_label(labels) # Convert labels to 0/1
    outputs = net(inputs)
    loss = criterion(outputs, labels.float())
    corr = (outputs > 0.0).squeeze().long() != labels
    total_err += int(corr.sum())
    total_loss += loss.item()
    total_epoch += len(labels)
err = float(total_err) / total_epoch
loss = float(total_loss) / (i + 1)
return err, loss
```

```
#####
```

**# Training Curve**

```
def plot_training_curve(path):
```

*""" Plots the training curve for a model run, given the csv files containing the train/validation error/loss.*

**Args:**

*path: The base path of the csv files produced during training*

```
"""
```

```
import matplotlib.pyplot as plt
train_err = np.loadtxt("{}_train_err.csv".format(path))
val_err = np.loadtxt("{}_val_err.csv".format(path))
train_loss = np.loadtxt("{}_train_loss.csv".format(path))
val_loss = np.loadtxt("{}_val_loss.csv".format(path))
plt.title("Train vs Validation Error")
n = len(train_err) # number of epochs
plt.plot(range(1,n+1), train_err, label="Train")
plt.plot(range(1,n+1), val_err, label="Validation")
plt.xlabel("Epoch")
plt.ylabel("Error")
plt.legend(loc='best')
plt.show()
plt.title("Train vs Validation Loss")
plt.plot(range(1,n+1), train_loss, label="Train")
plt.plot(range(1,n+1), val_loss, label="Validation")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend(loc='best')
plt.show()
```

## Part 1. Visualizing the Data [7 pt]

We will make use of some of the CIFAR-10 data set, which consists of colour images of size 32x32 pixels belonging to 10 categories. You can find out more about the dataset at

<https://www.cs.toronto.edu/~kriz/cifar.html> (<https://www.cs.toronto.edu/~kriz/cifar.html>)

For this assignment, we will only be using the cat and dog categories. We have included code that automatically downloads the dataset the first time that the main script is run.

In [ ]:

```
# This will download the CIFAR-10 dataset to a folder called "data"  
# the first time you run this code.  
train_loader, val_loader, test_loader, classes = get_data_loader(  
    target_classes=["cat", "dog"],  
    batch_size=1) # One image per batch
```

Files already downloaded and verified

Files already downloaded and verified

### Part (a) -- 1 pt

Visualize some of the data by running the code below. Include the visualization in your writeup.

(You don't need to submit anything else.)

In [ ]:


```

import matplotlib.pyplot as plt

k = 0
for images, labels in train_loader:
    # since batch_size = 1, there is only 1 image in `images`
    image = images[0]
    # place the colour channel at the end, instead of at the beginning
    img = np.transpose(image, [1,2,0])
    # normalize pixel intensity values to [0, 1]
    img = img / 2 + 0.5
    plt.subplot(3, 5, k+1)
    plt.axis('off')
    plt.imshow(img)

    k += 1
    if k > 14:
        break

```



## Part (b) -- 3 pt

How many training examples do we have for the combined `cat` and `dog` classes? What about validation examples? What about test examples?

In [ ]:

```
print("How many training examples do we have for the combined cat and dog classes?")
print("    " + str(len(train_loader)))

print("How many validation examples do we have for the combined cat and dog classes?")
print("    " + str(len(val_loader)))

print("How many training examples do we have for the combined cat and dog classes?")
print("    " + str(len(test_loader)))
```

How many training examples do we have for the combined cat and dog classes?

8000

How many validation examples do we have for the combined cat and dog classes?

2000

How many training examples do we have for the combined cat and dog classes?

2000

## Part (c) -- 3pt

Why do we need a validation set when training our model? What happens if we judge the performance of our models using the training set loss/error instead of the validation set loss/error?

Answer:

The validation set is used for testing models accuracy, parameter selection and to avoid overfitting.

The reason that we need a validation set when training our model is to make the model be more accurate. Since the testing datasets are for the final testing and should be used only when the model has been finished implemented, we should not use them while training.

If we use the training set, since we create the model by using them, if we want to make decisions about hyperparameters or track the accuracy, it will get overfit. The result will not be subjective to the adjustment and validation.

So, it is necessary to provide another set of validation data when training our model.

## Part 2. Training [15 pt]

We define two neural networks, a `LargeNet` and `SmallNet`. We'll be training the networks in this section.

You won't understand fully what these networks are doing until the next few classes, and that's okay. For this assignment, please focus on learning how to train networks, and how hyperparameters affect training.



In [ ]:

```

class LargeNet(nn.Module):
    def __init__(self):
        super(LargeNet, self).__init__()
        self.name = "large"
        self.conv1 = nn.Conv2d(3, 5, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(5, 10, 5)
        self.fc1 = nn.Linear(10 * 5 * 5, 32)
        self.fc2 = nn.Linear(32, 1)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 10 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        x = x.squeeze(1) # Flatten to [batch_size]
        return x

```

In [ ]:

```

class SmallNet(nn.Module):
    def __init__(self):
        super(SmallNet, self).__init__()
        self.name = "small"
        self.conv = nn.Conv2d(3, 5, 3)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc = nn.Linear(5 * 7 * 7, 1)

    def forward(self, x):
        x = self.pool(F.relu(self.conv(x)))
        x = self.pool(x)
        x = x.view(-1, 5 * 7 * 7)
        x = self.fc(x)
        x = x.squeeze(1) # Flatten to [batch_size]
        return x

```

In [ ]:

```

small_net = SmallNet()
large_net = LargeNet()

```

## Part (a) -- 2pt

The methods `small_net.parameters()` and `large_net.parameters()` produces an iterator of all the trainable parameters of the network. These parameters are torch tensors containing many scalar values.

We haven't learned how the parameters in these high-dimensional tensors will be used, but we should be able to count the number of parameters. Measuring the number of parameters in a network is one way of measuring the "size" of a network.

What is the total number of parameters in `small_net` and in `large_net` ? (Hint: how many numbers are in each tensor?)

In [ ]:

```
print("Small_net parameters:")
for param in small_net.parameters():
    print(param.shape)
print("\nLarge_net parameters:")
for param in large_net.parameters():
    print(param.shape)
```

Small\_net parameters:  
torch.Size([5, 3, 3, 3])  
torch.Size([5])  
torch.Size([1, 245])  
torch.Size([1])

Large\_net parameters:  
torch.Size([5, 3, 5, 5])  
torch.Size([5])  
torch.Size([10, 5, 5, 5])  
torch.Size([10])  
torch.Size([32, 250])  
torch.Size([32])  
torch.Size([1, 32])  
torch.Size([1])

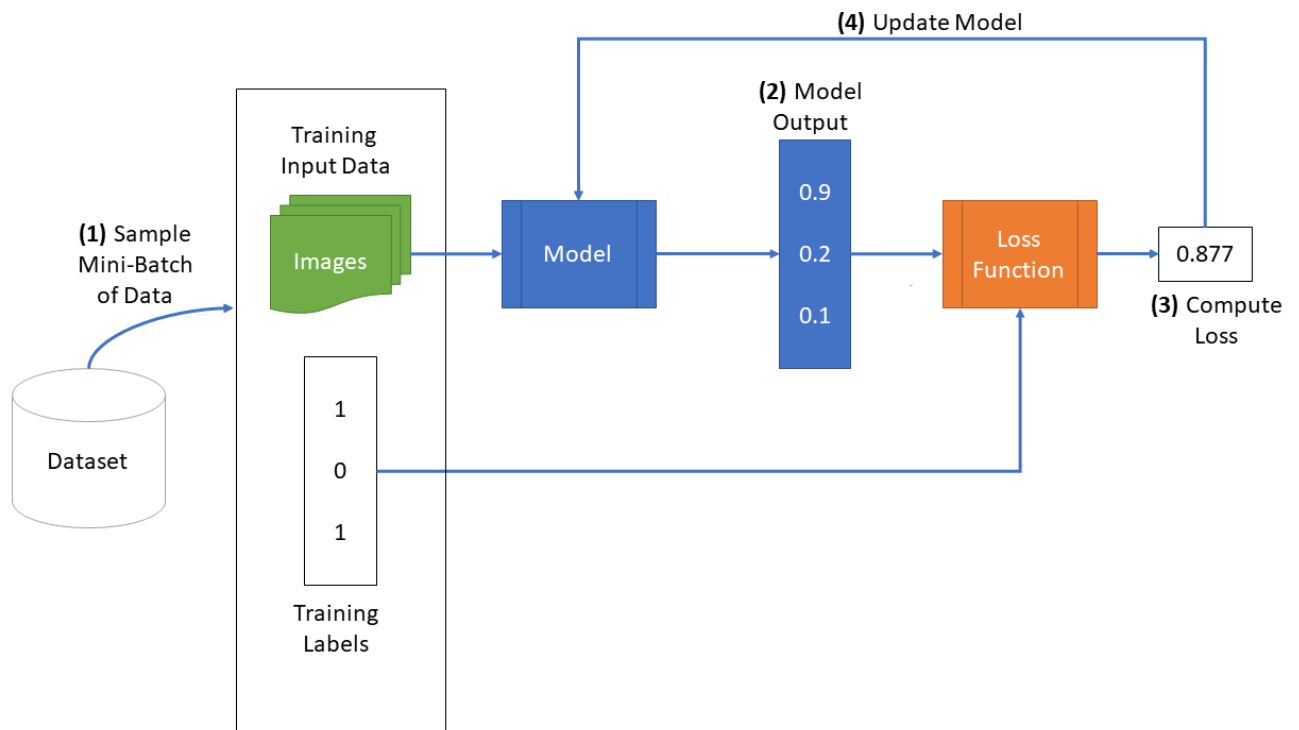
Answer:

Total # of parameters in small\_net =  $5333 + 5 + 1245 + 1 = 386$

Total # of parameters in large\_net =  $5355 + 5 + 10555 + 10 + 32250 + 32 + 132 + 1 = 9705$

## The function `train_net`

The function `train_net` below takes an untrained neural network (like `small_net` and `large_net`) and several other parameters. You should be able to understand how this function works. The figure below shows the high level training loop for a machine learning model:



In [ ]:

```

def train_net(net, batch_size=64, learning_rate=0.01, num_epochs=30):
    #####
    # Train a classifier on cats vs dogs
    target_classes = ["cat", "dog"]
    #####
    # Fixed PyTorch random seed for reproducible result
    torch.manual_seed(1000)
    #####
    # Obtain the PyTorch data loader objects to load batches of the datasets
    train_loader, val_loader, test_loader, classes = get_data_loader(
        target_classes, batch_size)
    #####
    # Define the Loss function and optimizer
    # The loss function will be Binary Cross Entropy (BCE). In this case we
    # will use the BCEWithLogitsLoss which takes unnormalized output from
    # the neural network and scalar label.
    # Optimizer will be SGD with Momentum.
    criterion = nn.BCEWithLogitsLoss()
    optimizer = optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)
    #####
    # Set up some numpy arrays to store the training/test loss/erruracy
    train_err = np.zeros(num_epochs)
    train_loss = np.zeros(num_epochs)
    val_err = np.zeros(num_epochs)
    val_loss = np.zeros(num_epochs)
    #####
    # Train the network
    # Loop over the data iterator and sample a new batch of training data
    # Get the output from the network, and optimize our loss function.
    start_time = time.time()
    for epoch in range(num_epochs): # loop over the dataset multiple times
        total_train_loss = 0.0
        total_train_err = 0.0
        total_epoch = 0
        for i, data in enumerate(train_loader, 0):
            # Get the inputs
            inputs, labels = data
            labels = normalize_label(labels) # Convert labels to 0/1
            # Zero the parameter gradients
            optimizer.zero_grad()
            # Forward pass, backward pass, and optimize
            outputs = net(inputs)
            loss = criterion(outputs, labels.float())
            loss.backward()
            optimizer.step()
            # Calculate the statistics
            corr = (outputs > 0.0).squeeze().long() != labels
            total_train_err += int(corr.sum())
            total_train_loss += loss.item()
            total_epoch += len(labels)
        train_err[epoch] = float(total_train_err) / total_epoch
        train_loss[epoch] = float(total_train_loss) / (i+1)
        val_err[epoch], val_loss[epoch] = evaluate(net, val_loader, criterion)
        print(("Epoch {}: Train err: {}, Train loss: {} | "+
            "Validation err: {}, Validation loss: {}".format(
                epoch + 1,
                train_err[epoch],
                train_loss[epoch],
                val_err[epoch],

```

```

        val_loss[epoch]))
    # Save the current model (checkpoint) to a file
    model_path = get_model_name(net.name, batch_size, learning_rate, epoch)
    torch.save(net.state_dict(), model_path)
print('Finished Training')
end_time = time.time()
elapsed_time = end_time - start_time
print("Total time elapsed: {:.2f} seconds".format(elapsed_time))
# Write the train/test loss/err into CSV file for plotting later
epochs = np.arange(1, num_epochs + 1)
np.savetxt("{}_train_err.csv".format(model_path), train_err)
np.savetxt("{}_train_loss.csv".format(model_path), train_loss)
np.savetxt("{}_val_err.csv".format(model_path), val_err)
np.savetxt("{}_val_loss.csv".format(model_path), val_loss)

```

## Part (b) -- 1pt

The parameters to the function `train_net` are hyperparameters of our neural network. We made these hyperparameters easy to modify so that we can tune them later on.

What are the default values of the parameters `batch_size`, `learning_rate`, and `num_epochs`?

Answer:

The default values of the parameters `batch_size`, `learning_rate`, and `num_epochs` are 64, 0.01 and 30, respectively.

## Part (c) -- 3 pt

What files are written to disk when we call `train_net` with `small_net`, and train for 5 epochs? Provide a list of all the files written to disk, and what information the files contain.

In [ ]:

```

train_net(small_net, num_epochs=5) #call train_net with small_net, and train for 5 epochs

```

```

Files already downloaded and verified
Files already downloaded and verified
Epoch 1: Train err: 0.429125, Train loss: 0.6759076609611511 |Validation err: 0.3775, Validation loss: 0.6534486319869757
Epoch 2: Train err: 0.361, Train loss: 0.6403536524772644 |Validation err: 0.3775, Validation loss: 0.6540046762675047
Epoch 3: Train err: 0.334625, Train loss: 0.6169258975982665 |Validation err: 0.336, Validation loss: 0.6135102435946465
Epoch 4: Train err: 0.318875, Train loss: 0.5990912899971008 |Validation err: 0.3485, Validation loss: 0.6191740091890097
Epoch 5: Train err: 0.309875, Train loss: 0.5884846107959747 |Validation err: 0.3285, Validation loss: 0.6017458196729422
Finished Training
Total time elapsed: 19.58 seconds

```

Answer:

Based on the following code, there exists five checkpoints that it saves the current model to a file.

```
model_path = get_model_name(net.name, batch_size, learning_rate, epoch)
torch.save(net.state_dict(), model_path)
```

Since we are calling the `small_net` and no change with the `batch_size` and `learning_rate`, the five file saving process will be:

1. Model: `small_net`, Batch\_size: 64, Learning\_rate: 0.01, Epoch: 0
2. Model: `small_net`, Batch\_size: 64, Learning\_rate: 0.01, Epoch: 1
3. Model: `small_net`, Batch\_size: 64, Learning\_rate: 0.01, Epoch: 2
4. Model: `small_net`, Batch\_size: 64, Learning\_rate: 0.01, Epoch: 3
5. Model: `small_net`, Batch\_size: 64, Learning\_rate: 0.01, Epoch: 4

Based on the following code, there exists four times of saving files to CSV type for the train/test loss/err values.

```
epochs = np.arange(1, num_epochs + 1)

np.savetxt("{}_train_err.csv".format(model_path), train_err)

np.savetxt("{}_train_loss.csv".format(model_path), train_loss)

np.savetxt("{}_val_err.csv".format(model_path), val_err)

np.savetxt("{}_val_loss.csv".format(model_path), val_loss)
```

Since the code is outside the epoch for loop, it saved the `train_err`, `train_loss`, `val_err` and `val_loss` datas separately with four times in total.

1. Model: `small_net`, Batch\_size: 64, Learning\_rate: 0.01, Epoch: 4, CSV file content: `train_err`
2. Model: `small_net`, Batch\_size: 64, Learning\_rate: 0.01, Epoch: 4, CSV file content: `train_loss`
3. Model: `small_net`, Batch\_size: 64, Learning\_rate: 0.01, Epoch: 4, CSV file content: `val_err`
4. Model: `small_net`, Batch\_size: 64, Learning\_rate: 0.01, Epoch: 4, CSV file content: `val_loss`

## Part (d) -- 2pt

Train both `small_net` and `large_net` using the function `train_net` and its default parameters. The function will write many files to disk, including a model checkpoint (saved values of model weights) at the end of each epoch.

If you are using Google Colab, you will need to mount Google Drive so that the files generated by `train_net` gets saved. We will be using these files in part (d). (See the Google Colab tutorial for more information about this.)

Report the total time elapsed when training each network. Which network took longer to train? Why?

In [ ]:

```
# Since the function writes files to disk, you will need to mount  
# your Google Drive. If you are working on the lab locally, you  
# can comment out this code.
```

```
from google.colab import drive  
drive.mount('/content/gdrive')
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force\_remount=True).

In [ ]:

```
# Train both small_net and large_net using the function train_net and its default parameters
print("-----training for small_net")
train_net(small_net)

print("\n-----training for large_net")
train_net(large_net)
```



```
-----training for small_n
et
Files already downloaded and verified
Files already downloaded and verified
Epoch 1: Train err: 0.426375, Train loss: 0.6762333779335022 |Valida
tion err: 0.369, Validation loss: 0.6512321196496487
Epoch 2: Train err: 0.354875, Train loss: 0.6383493151664734 |Valida
tion err: 0.3645, Validation loss: 0.6542029641568661
Epoch 3: Train err: 0.3395, Train loss: 0.6214949712753296 |Validati
on err: 0.35, Validation loss: 0.6223773881793022
Epoch 4: Train err: 0.3315, Train loss: 0.6070578649044037 |Validati
on err: 0.3495, Validation loss: 0.6307939775288105
Epoch 5: Train err: 0.320625, Train loss: 0.5986896252632141 |Valida
tion err: 0.3315, Validation loss: 0.6189757902175188
Epoch 6: Train err: 0.30975, Train loss: 0.5889757220745087 |Validat
ion err: 0.324, Validation loss: 0.6137909144163132
Epoch 7: Train err: 0.30975, Train loss: 0.5846510627269745 |Validat
ion err: 0.332, Validation loss: 0.6105258176103234
Epoch 8: Train err: 0.3025, Train loss: 0.5786999728679657 |Validati
on err: 0.3245, Validation loss: 0.6059290692210197
Epoch 9: Train err: 0.295375, Train loss: 0.5760825304985047 |Valida
tion err: 0.322, Validation loss: 0.6085914196446538
Epoch 10: Train err: 0.289875, Train loss: 0.5690126593112945 |Valid
ation err: 0.313, Validation loss: 0.6022631283849478
Epoch 11: Train err: 0.29125, Train loss: 0.5676470882892609 |Valida
tion err: 0.31, Validation loss: 0.5932684279978275
Epoch 12: Train err: 0.290875, Train loss: 0.5611145758628845 |Valid
ation err: 0.3225, Validation loss: 0.6004137145355344
Epoch 13: Train err: 0.287625, Train loss: 0.5621939980983734 |Valid
ation err: 0.302, Validation loss: 0.5935621615499258
Epoch 14: Train err: 0.28525, Train loss: 0.5567016141414642 |Valida
tion err: 0.322, Validation loss: 0.6071959668770432
Epoch 15: Train err: 0.283, Train loss: 0.553466688156128 |Validatio
n err: 0.308, Validation loss: 0.5929243136197329
Epoch 16: Train err: 0.283875, Train loss: 0.5557780301570893 |Valid
ation err: 0.305, Validation loss: 0.5938641577959061
Epoch 17: Train err: 0.275, Train loss: 0.5508679287433624 |Validati
on err: 0.299, Validation loss: 0.5814803345128894
Epoch 18: Train err: 0.279625, Train loss: 0.5483031418323517 |Valid
ation err: 0.299, Validation loss: 0.5813600467517972
Epoch 19: Train err: 0.274375, Train loss: 0.5433929843902587 |Valid
ation err: 0.312, Validation loss: 0.5894149150699377
Epoch 20: Train err: 0.271375, Train loss: 0.5439429087638855 |Valid
ation err: 0.298, Validation loss: 0.5814473712816834
Epoch 21: Train err: 0.273625, Train loss: 0.5445237925052643 |Valid
ation err: 0.301, Validation loss: 0.5790077913552523
Epoch 22: Train err: 0.2715, Train loss: 0.5407345817089081 |Validat
ion err: 0.31, Validation loss: 0.5918727135285735
Epoch 23: Train err: 0.271375, Train loss: 0.5413676164150238 |Valid
ation err: 0.299, Validation loss: 0.5804261956363916
Epoch 24: Train err: 0.270625, Train loss: 0.5376549961566925 |Valid
ation err: 0.303, Validation loss: 0.5810985947027802
Epoch 25: Train err: 0.269875, Train loss: 0.5345538351535797 |Valid
ation err: 0.306, Validation loss: 0.5838705860078335
Epoch 26: Train err: 0.26825, Train loss: 0.5351638784408569 |Valida
tion err: 0.292, Validation loss: 0.5764822503551841
Epoch 27: Train err: 0.266375, Train loss: 0.5333331351280213 |Valid
ation err: 0.3005, Validation loss: 0.5876168962568045
Epoch 28: Train err: 0.270125, Train loss: 0.5356989531517029 |Valid
ation err: 0.29, Validation loss: 0.5697013400495052
Epoch 29: Train err: 0.266125, Train loss: 0.5333648209571838 |Valid
```

ation err: 0.2965, Validation loss: 0.5893930019810796  
Epoch 30: Train err: 0.27125, Train loss: 0.53548370885849 |Validation  
on err: 0.307, Validation loss: 0.5843431651592255  
Finished Training  
Total time elapsed: 112.14 seconds

-----training for large\_n  
et

Files already downloaded and verified

Files already downloaded and verified

Epoch 1: Train err: 0.45175, Train loss: 0.6888214983940124 |Validation  
err: 0.445, Validation loss: 0.6789547484368086

Epoch 2: Train err: 0.41575, Train loss: 0.674795910358429 |Validation  
err: 0.413, Validation loss: 0.672644479200244

Epoch 3: Train err: 0.389375, Train loss: 0.6573688311576843 |Valida  
tion err: 0.378, Validation loss: 0.6426923684775829

Epoch 4: Train err: 0.361, Train loss: 0.634773886680603 |Validation  
err: 0.3505, Validation loss: 0.6394727751612663

Epoch 5: Train err: 0.341375, Train loss: 0.6152373743057251 |Valida  
tion err: 0.3355, Validation loss: 0.616610599681735

Epoch 6: Train err: 0.325, Train loss: 0.6017866497039794 |Validatio  
n err: 0.3315, Validation loss: 0.6125319022685289

Epoch 7: Train err: 0.32, Train loss: 0.5897287225723267 |Validation  
err: 0.342, Validation loss: 0.611613916233182

Epoch 8: Train err: 0.31025, Train loss: 0.5780349550247192 |Validat  
ion err: 0.3165, Validation loss: 0.5924079436808825

Epoch 9: Train err: 0.298875, Train loss: 0.5735508439540863 |Valida  
tion err: 0.3255, Validation loss: 0.5936816073954105

Epoch 10: Train err: 0.289375, Train loss: 0.558117666721344 |Valida  
tion err: 0.3095, Validation loss: 0.5808845506981015

Epoch 11: Train err: 0.281625, Train loss: 0.5475617473125458 |Valid  
ation err: 0.3095, Validation loss: 0.5984619250521064

Epoch 12: Train err: 0.27075, Train loss: 0.5367428750991822 |Valida  
tion err: 0.2995, Validation loss: 0.591126230545342

Epoch 13: Train err: 0.262375, Train loss: 0.525796151638031 |Valida  
tion err: 0.299, Validation loss: 0.5876793833449483

Epoch 14: Train err: 0.2615, Train loss: 0.5151971943378448 |Validat  
ion err: 0.298, Validation loss: 0.5926839364692569

Epoch 15: Train err: 0.252375, Train loss: 0.5053955318927765 |Valid  
ation err: 0.299, Validation loss: 0.5844418630003929

Epoch 16: Train err: 0.245125, Train loss: 0.4993929398059845 |Valid  
ation err: 0.2965, Validation loss: 0.5912785157561302

Epoch 17: Train err: 0.236625, Train loss: 0.4863951737880707 |Valid  
ation err: 0.309, Validation loss: 0.5879802675917745

Epoch 18: Train err: 0.228, Train loss: 0.47127222537994384 |Validat  
ion err: 0.2945, Validation loss: 0.5824780873954296

Epoch 19: Train err: 0.223375, Train loss: 0.4614680805206299 |Valid  
ation err: 0.295, Validation loss: 0.5957808950915933

Epoch 20: Train err: 0.216125, Train loss: 0.4512100396156311 |Valid  
ation err: 0.322, Validation loss: 0.6553378133103251

Epoch 21: Train err: 0.217625, Train loss: 0.4489641053676605 |Valid  
ation err: 0.296, Validation loss: 0.5930969789624214

Epoch 22: Train err: 0.202625, Train loss: 0.43051408672332764 |Vali  
dation err: 0.3165, Validation loss: 0.6367753250524402

Epoch 23: Train err: 0.19825, Train loss: 0.41719653725624084 |Valid  
ation err: 0.2975, Validation loss: 0.5902507230639458

Epoch 24: Train err: 0.194375, Train loss: 0.41712852954864504 |Vali  
dation err: 0.2925, Validation loss: 0.6170288324356079

Epoch 25: Train err: 0.17725, Train loss: 0.391695091843605 |Validat  
ion err: 0.304, Validation loss: 0.6329760905355215

Epoch 26: Train err: 0.169125, Train loss: 0.3703381236791611 |Valid

```
ation err: 0.3305, Validation loss: 0.6798105845227838
Epoch 27: Train err: 0.164875, Train loss: 0.36250311982631683 |Vali
dation err: 0.3095, Validation loss: 0.6954353488981724
Epoch 28: Train err: 0.160375, Train loss: 0.35222707903385164 |Vali
dation err: 0.3185, Validation loss: 0.710120321251452
Epoch 29: Train err: 0.1435, Train loss: 0.3215026478767395 |Validat
ion err: 0.321, Validation loss: 0.8275542985647917
Epoch 30: Train err: 0.140125, Train loss: 0.3160476162433624 |Valid
ation err: 0.317, Validation loss: 0.7528756745159626
Finished Training
Total time elapsed: 124.81 seconds
```

Answer:

According to the output above, the training time elapsed for `small_net` process and `large_net` process are 98.00 seconds and 106.67 seconds , respectively.

The `large_net` network takes more time, because it contains more parameters that the `small_net` network and it takes more time to modeling the training process.

I need to state that, the first time I run the test cases, the output are 98.00 seconds and 106.67 seconds. I kept this two number for comparison for the following part in part2 and part3. However, after I rerun the code as preperation for the later parts (4&5) the other day. The time elapsed has changed to 115.75 seconds and 128.97 seconds. Each time I ran, it has small variations. But, overall, the `large_net` network will take more time than the `small_net` network.

## Part (e) - 2pt

Use the function `plot_training_curve` to display the trajectory of the training/validation error and the training/validation loss. You will need to use the function `get_model_name` to generate the argument to the `plot_training_curve` function.


Do this for both the small network and the large network. Include both plots in your writeup.

In [ ]:



```
print("-----output curve for small_net")
small_model_path = get_model_name("small", batch_size=64, learning_rate=0.01, epoch=29)
plot_training_curve(small_model_path)

print("\n-----output curve for large_net")
large_model_path = get_model_name("large", batch_size=64, learning_rate=0.01, epoch=29)
plot_training_curve(large_model_path)
```

-----output curve for sma  
ll\_net



-----output curve for lar  
ge\_net



## Part (f) - 5pt

Describe what you notice about the training curve. How do the curves differ for `small_net` and `large_net` ? Identify any occurrences of underfitting and overfitting.

According to the lecture, overfitting occurs when training data performs well but validation data performs poorly. Underfitting is when the model is not complex enough for all relationships for inputs and labels, so it has high bias.

As we noticed from the `small_net` network, the general trends as the epoch number increases, the training and validation data decreases are similar. The difference between error/loss of the training and validation data are not large.

For the `large_net` network, we realized that at the beginning the loss and error values are significant, which implies that the model is underfitting. After more epoch iterations had occurred, the training datasets reached a very low error and loss, but leads to a very high or even higher validation error. So, it turns out to have a sign of overfitting, which is different from the `small_net` network in the general trends.

## Part 3. Optimization Parameters [12 pt]

For this section, we will work with `large_net` only.

### Part (a) - 3pt

Train `large_net` with all default parameters, except set `learning_rate=0.001`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *lowering* the learning rate.

In [ ]:

```
# Note: When we re-construct the model, we start the training  
# with *random weights*. If we omit this code, the values of  
# the weights will still be the previously trained values.  
large_net = LargeNet()  
train_net(large_net, 64, 0.001, 30) #Train large_net with all default parameter  
s, except set learning_rate=0.001  
  
print("\n-----output curve for large_  
net")  
large_model_path = get_model_name("large", batch_size=64, learning_rate=0.001, e  
poch=29)  
plot_training_curve(large_model_path)
```




Files already downloaded and verified

Files already downloaded and verified

Epoch 1: Train err: 0.47625, Train loss: 0.6928360018730163 | Validation err: 0.467, Validation loss: 0.6924686655402184  
Epoch 2: Train err: 0.448625, Train loss: 0.692258969783783 | Validation err: 0.4305, Validation loss: 0.691649341955781  
Epoch 3: Train err: 0.43575, Train loss: 0.6916067352294922 | Validation err: 0.4285, Validation loss: 0.6908544301986694  
Epoch 4: Train err: 0.43, Train loss: 0.6908614501953125 | Validation err: 0.4245, Validation loss: 0.6896599512547255  
Epoch 5: Train err: 0.434375, Train loss: 0.6899197044372558 | Validation err: 0.4195, Validation loss: 0.6886946316808462  
Epoch 6: Train err: 0.435875, Train loss: 0.6887416214942932 | Validation err: 0.4195, Validation loss: 0.6867832522839308  
Epoch 7: Train err: 0.436875, Train loss: 0.687377580165863 | Validation err: 0.418, Validation loss: 0.6851987168192863  
Epoch 8: Train err: 0.4375, Train loss: 0.6859274730682373 | Validation err: 0.4115, Validation loss: 0.6831987891346216  
Epoch 9: Train err: 0.424625, Train loss: 0.6844045362472534 | Validation err: 0.412, Validation loss: 0.6808866523206234  
Epoch 10: Train err: 0.424, Train loss: 0.6828485646247864 | Validation err: 0.408, Validation loss: 0.6783493831753731  
Epoch 11: Train err: 0.42525, Train loss: 0.6812345442771912 | Validation err: 0.4125, Validation loss: 0.6780271530151367  
Epoch 12: Train err: 0.42, Train loss: 0.6796339163780213 | Validation err: 0.4125, Validation loss: 0.6753249019384384  
Epoch 13: Train err: 0.41475, Train loss: 0.6777957553863525 | Validation err: 0.415, Validation loss: 0.6757064927369356  
Epoch 14: Train err: 0.412125, Train loss: 0.6761115918159485 | Validation err: 0.412, Validation loss: 0.6739722862839699  
Epoch 15: Train err: 0.40925, Train loss: 0.6744697222709656 | Validation err: 0.4145, Validation loss: 0.6706754509359598  
Epoch 16: Train err: 0.406125, Train loss: 0.6727363333702088 | Validation err: 0.4105, Validation loss: 0.6707725338637829  
Epoch 17: Train err: 0.401375, Train loss: 0.6713055086135864 | Validation err: 0.4035, Validation loss: 0.6671533733606339  
Epoch 18: Train err: 0.399125, Train loss: 0.6696757755279541 | Validation err: 0.4055, Validation loss: 0.6646822281181812  
Epoch 19: Train err: 0.401, Train loss: 0.6679071426391602 | Validation err: 0.396, Validation loss: 0.6655134689062834  
Epoch 20: Train err: 0.392125, Train loss: 0.6657885494232177 | Validation err: 0.405, Validation loss: 0.6626032534986734  
Epoch 21: Train err: 0.38975, Train loss: 0.6646304850578308 | Validation err: 0.3945, Validation loss: 0.6606811657547951  
Epoch 22: Train err: 0.388625, Train loss: 0.6623743529319763 | Validation err: 0.3935, Validation loss: 0.6617076825350523  
Epoch 23: Train err: 0.384125, Train loss: 0.6601564373970031 | Validation err: 0.3965, Validation loss: 0.6574041564017534  
Epoch 24: Train err: 0.382125, Train loss: 0.6584046039581298 | Validation err: 0.3855, Validation loss: 0.656140647828579  
Epoch 25: Train err: 0.378375, Train loss: 0.6555113220214843 | Validation err: 0.388, Validation loss: 0.6552989799529314  
Epoch 26: Train err: 0.376875, Train loss: 0.6531331434249877 | Validation err: 0.3875, Validation loss: 0.6531767845153809  
Epoch 27: Train err: 0.375, Train loss: 0.6503931031227111 | Validation err: 0.387, Validation loss: 0.6520250719040632  
Epoch 28: Train err: 0.37125, Train loss: 0.6476753010749817 | Validation err: 0.3875, Validation loss: 0.6483868304640055  
Epoch 29: Train err: 0.367875, Train loss: 0.6451566920280457 | Validation err: 0.3815, Validation loss: 0.6459947340190411  
Epoch 30: Train err: 0.362375, Train loss: 0.6423771266937256 | Validation err: 0.3815, Validation loss: 0.6459947340190411

ation err: 0.379, Validation loss: 0.6439598128199577  
Finished Training  
Total time elapsed: 106.52 seconds

-----output curve for lar  
ge\_net



Answer:

With `learning_rate = 0.001`, it takes 106.52 seconds. With `learning_rate = 0.01`, it takes 106.67 seconds. These two testcases roughly takes the same time for modeling.

---

We noticed by analysing the graph, by lowering the `learning_rate`, the training and validation value are roughly have similar values, which represents it can avoid overfitting compares to the previous diagram with smaller step size for the weights updated during the training process.

Also, for the previous error/loss diagram, when epoch closer to 29 and learning rate = 0.01, the error/loss value are nearly 0.15 and 0.3. However, in this case, when epoch closer to 29 and learning rate = 0.001, the error/loss value are significantly higher than 0.36 and 0.64. The errorness and loss increased to be higher than before, which means 0.001 is still not be the best case as when updating the weights, the step size is too small, so it is too slow within the 30 epoches to reach a very good model/weights. According to lecture notes, for smaller learning rates, parameters don't change very much in each iteration and it will take a long time to train the network.

## Part (b) - 3pt

Train `large_net` with all default parameters, except set `learning_rate=0.1`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *increasing* the learning rate.

In [ ]:

```
large_net = LargeNet()
train_net(large_net, 64, 0.1, 30) #Train large_net with all default parameters,
except set learning_rate=0.1

print("\n-----output curve for large_
net")
large_model_path = get_model_name("large", batch_size=64, learning_rate=0.1, epo
ch=29)
plot_training_curve(large_model_path)
```

Files already downloaded and verified

Files already downloaded and verified

Epoch 1: Train err: 0.426875, Train loss: 0.6743021264076233 | Validation err: 0.373, Validation loss: 0.6364669129252434  
Epoch 2: Train err: 0.367125, Train loss: 0.6398710680007934 | Validation err: 0.3755, Validation loss: 0.6349833384156227  
Epoch 3: Train err: 0.358125, Train loss: 0.629592146396637 | Validation err: 0.348, Validation loss: 0.6240996662527323  
Epoch 4: Train err: 0.343125, Train loss: 0.6167744815349578 | Validation err: 0.3405, Validation loss: 0.6115005295723677  
Epoch 5: Train err: 0.3315, Train loss: 0.5990285336971283 | Validation err: 0.348, Validation loss: 0.6214024610817432  
Epoch 6: Train err: 0.31125, Train loss: 0.5800493450164795 | Validation err: 0.3195, Validation loss: 0.6157716857269406  
Epoch 7: Train err: 0.31025, Train loss: 0.5827681159973145 | Validation err: 0.3625, Validation loss: 0.6364191770553589  
Epoch 8: Train err: 0.305375, Train loss: 0.5704554648399353 | Validation err: 0.3365, Validation loss: 0.6044664867222309  
Epoch 9: Train err: 0.290125, Train loss: 0.5561789710521698 | Validation err: 0.326, Validation loss: 0.5908120637759566  
Epoch 10: Train err: 0.280125, Train loss: 0.5481600475311279 | Validation err: 0.335, Validation loss: 0.6053934590891004  
Epoch 11: Train err: 0.279125, Train loss: 0.5420139644145966 | Validation err: 0.325, Validation loss: 0.6179495453834534  
Epoch 12: Train err: 0.26675, Train loss: 0.5261815791130066 | Validation err: 0.314, Validation loss: 0.6270146453753114  
Epoch 13: Train err: 0.261875, Train loss: 0.5207161004543305 | Validation err: 0.3415, Validation loss: 0.6598431821912527  
Epoch 14: Train err: 0.27, Train loss: 0.5271751899719238 | Validation err: 0.3355, Validation loss: 0.6816458441317081  
Epoch 15: Train err: 0.25375, Train loss: 0.5104072098731994 | Validation err: 0.3325, Validation loss: 0.7142039136961102  
Epoch 16: Train err: 0.252625, Train loss: 0.5123440861701966 | Validation err: 0.337, Validation loss: 0.668152442201972  
Epoch 17: Train err: 0.241375, Train loss: 0.5032827985286713 | Validation err: 0.336, Validation loss: 0.6614612024277449  
Epoch 18: Train err: 0.228125, Train loss: 0.4766872684955597 | Validation err: 0.351, Validation loss: 0.6638622563332319  
Epoch 19: Train err: 0.2355, Train loss: 0.4919496657848358 | Validation err: 0.3615, Validation loss: 0.7392734382301569  
Epoch 20: Train err: 0.24525, Train loss: 0.4942772650718689 | Validation err: 0.33, Validation loss: 0.701873391866684  
Epoch 21: Train err: 0.2415, Train loss: 0.49403335285186767 | Validation err: 0.3475, Validation loss: 0.7289880393072963  
Epoch 22: Train err: 0.236125, Train loss: 0.48673622250556947 | Validation err: 0.3405, Validation loss: 0.8085937779396772  
Epoch 23: Train err: 0.220375, Train loss: 0.4598086004257202 | Validation err: 0.3195, Validation loss: 0.7529164329171181  
Epoch 24: Train err: 0.22675, Train loss: 0.47790721583366397 | Validation err: 0.3505, Validation loss: 0.8729726411402225  
Epoch 25: Train err: 0.22825, Train loss: 0.473807953119278 | Validation err: 0.365, Validation loss: 0.7805169755592942  
Epoch 26: Train err: 0.21425, Train loss: 0.45517154574394225 | Validation err: 0.3485, Validation loss: 0.7850301470607519  
Epoch 27: Train err: 0.220875, Train loss: 0.47447798085212706 | Validation err: 0.34, Validation loss: 0.7402617856860161  
Epoch 28: Train err: 0.214125, Train loss: 0.454453275680542 | Validation err: 0.334, Validation loss: 0.7570528583601117  
Epoch 29: Train err: 0.222, Train loss: 0.47537617254257203 | Validation err: 0.345, Validation loss: 1.017714038491249  
Epoch 30: Train err: 0.228875, Train loss: 0.4917977278232574 | Validation err: 0.345, Validation loss: 0.7805169755592942

ation err: 0.354, Validation loss: 0.9331185091286898  
Finished Training  
Total time elapsed: 106.71 seconds

-----output curve for lar  
ge\_net



Answer:

With `learning_rate = 0.1`, it takes 106.71 seconds. With `learning_rate = 0.01`, it takes 106.67 seconds. These two testcases roughly takes the same time for modeling.

---

We noticed by analysing the graph, by increasing the `learning_rate`, the training and validation value have larger separation for the trends, which represents it cannot avoid overfitting and enlarged the overfitting issue compares to the previous diagram, with larger step size for the weights updated during the training process.

Also, for the previous error/loss diagram, when epoch closer to 29 and learning rate = 0.01, the error/loss value are nearly 0.15 and 0.3. However, in this case, when epoch closer to 29 and learning rate = 0.1, the error/loss value are significantly higher than 0.36 and 0.9. The errorness and loss increased to be higher than before, especially for validation datasets, although the training datasets has a lower loss/err value. It means 0.1 is still not be the best case as when updating the weights, the step size is too big, it is too fast within the 30 epoches and enlarge the overfitting of the model. According to the lecture notes, large learning rate has large step size that can be detrimental to neural network training.

### Part (c) - 3pt

Train `large_net` with all default parameters, including with `learning_rate=0.01`. Now, set `batch_size=512`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *increasing* the batch size.

In [ ]:

```
large_net = LargeNet()
train_net(large_net, 512, 0.01, 30) #Now, set batch_size=512

print("\n-----output curve for large_
net")
large_model_path = get_model_name("large", batch_size=512, learning_rate=0.01, e
poch=29)
plot_training_curve(large_model_path)
```




Files already downloaded and verified

Files already downloaded and verified

Epoch 1: Train err: 0.48175, Train loss: 0.6929379515349865 | Validation err: 0.478, Validation loss: 0.6926823854446411  
Epoch 2: Train err: 0.457625, Train loss: 0.6924103982746601 | Validation err: 0.434, Validation loss: 0.6917425245046616  
Epoch 3: Train err: 0.437, Train loss: 0.6916500590741634 | Validation err: 0.4265, Validation loss: 0.6909129917621613  
Epoch 4: Train err: 0.433625, Train loss: 0.6908449940383434 | Validation err: 0.424, Validation loss: 0.6897870451211929  
Epoch 5: Train err: 0.434, Train loss: 0.6896935552358627 | Validation err: 0.424, Validation loss: 0.6881355047225952  
Epoch 6: Train err: 0.438, Train loss: 0.6883532032370567 | Validation err: 0.4285, Validation loss: 0.686011865735054  
Epoch 7: Train err: 0.439375, Train loss: 0.6866871826350689 | Validation err: 0.426, Validation loss: 0.6836968809366226  
Epoch 8: Train err: 0.43525, Train loss: 0.6849770694971085 | Validation err: 0.411, Validation loss: 0.6814672648906708  
Epoch 9: Train err: 0.42375, Train loss: 0.6832008808851242 | Validation err: 0.414, Validation loss: 0.6795913726091385  
Epoch 10: Train err: 0.421, Train loss: 0.6811088100075722 | Validation err: 0.416, Validation loss: 0.677154153585434  
Epoch 11: Train err: 0.42075, Train loss: 0.6794030219316483 | Validation err: 0.4095, Validation loss: 0.6748131364583969  
Epoch 12: Train err: 0.41475, Train loss: 0.6768063381314278 | Validation err: 0.412, Validation loss: 0.673705518245697  
Epoch 13: Train err: 0.41025, Train loss: 0.6749699972569942 | Validation err: 0.412, Validation loss: 0.6706137955188751  
Epoch 14: Train err: 0.40725, Train loss: 0.6730927303433418 | Validation err: 0.4125, Validation loss: 0.6692113280296326  
Epoch 15: Train err: 0.400375, Train loss: 0.6706851311028004 | Validation err: 0.4105, Validation loss: 0.6672530621290207  
Epoch 16: Train err: 0.397625, Train loss: 0.6691787466406822 | Validation err: 0.405, Validation loss: 0.6649062037467957  
Epoch 17: Train err: 0.394, Train loss: 0.6675723902881145 | Validation err: 0.401, Validation loss: 0.6630214899778366  
Epoch 18: Train err: 0.393, Train loss: 0.6648031920194626 | Validation err: 0.394, Validation loss: 0.6623998433351517  
Epoch 19: Train err: 0.386375, Train loss: 0.6627435684204102 | Validation err: 0.3875, Validation loss: 0.6597269624471664  
Epoch 20: Train err: 0.381875, Train loss: 0.6596166752278805 | Validation err: 0.4005, Validation loss: 0.6564314365386963  
Epoch 21: Train err: 0.38575, Train loss: 0.6584814041852951 | Validation err: 0.388, Validation loss: 0.6586655974388123  
Epoch 22: Train err: 0.378375, Train loss: 0.6551217511296272 | Validation err: 0.385, Validation loss: 0.652855172753334  
Epoch 23: Train err: 0.3715, Train loss: 0.6508790552616119 | Validation err: 0.3835, Validation loss: 0.6497973203659058  
Epoch 24: Train err: 0.376625, Train loss: 0.6488082818686962 | Validation err: 0.3845, Validation loss: 0.6474854201078415  
Epoch 25: Train err: 0.369125, Train loss: 0.6445849314332008 | Validation err: 0.3815, Validation loss: 0.647366538643837  
Epoch 26: Train err: 0.373, Train loss: 0.6428744271397591 | Validation err: 0.3745, Validation loss: 0.6425762921571732  
Epoch 27: Train err: 0.359375, Train loss: 0.6372172273695469 | Validation err: 0.378, Validation loss: 0.6397283673286438  
Epoch 28: Train err: 0.35425, Train loss: 0.6337856389582157 | Validation err: 0.3705, Validation loss: 0.6403937339782715  
Epoch 29: Train err: 0.35375, Train loss: 0.6311305873095989 | Validation err: 0.3665, Validation loss: 0.6336003690958023  
Epoch 30: Train err: 0.352625, Train loss: 0.6283440999686718 | Validation err: 0.3665, Validation loss: 0.6336003690958023

ation err: 0.3665, Validation loss: 0.6324431151151657  
Finished Training  
Total time elapsed: 93.42 seconds

-----output curve for lar  
ge\_net



Answer:

With `batch_size = 512`, it takes 93.42 seconds. With `batch_size = 64`, it takes 106.67 seconds. With higher `batch_size`, the network takes less time to complete modeling.

As `batch_size` is representing the number of samples per batch. So, if by increasing the `batch_size`, which means that the number of samples per batch will increase and the number of batches itself decreases. So, it will take less time for each iteration and leads to an overall less time.

---

We noticed by analysing the graph, by increasing the `batch_size`, the training and validation value lines are roughly have similar values, which represents it can avoid overfitting compared to the previous diagram.

Also, for the previous error/loss diagram, when epoch closer to 29 and `batch_size = 64`, the error/loss value are nearly 0.15 and 0.3. However, in this case, when epoch closer to 29 and `batch_size = 512`, the error/loss value are significantly higher than 0.36 and 0.63. The error and loss increased to be higher than before, as larger `batch_size` represents for the larger amount of samples that has been calculated each iteration for the gradient and direction, which leads to the incorrectness. According to the lecture notes, since average loss might not change very much as batch size grows, so larger batch size is too precise as it limits the "noise" occur in the iteration and let the training converge slower.

## Part (d) - 3pt

Train `large_net` with all default parameters, including with `learning_rate=0.01`. Now, set `batch_size=16`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *decreasing* the batch size.

In [ ]:

```
large_net = LargeNet()
train_net(large_net, 16, 0.01, 30) #Now, set batch_size=16

print("\n-----output curve for large_
net")
large_model_path = get_model_name("large", batch_size=16, learning_rate=0.01, epoch=29)
plot_training_curve(large_model_path)
```

Files already downloaded and verified

Files already downloaded and verified

Epoch 1: Train err: 0.432, Train loss: 0.6775538735389709 | Validation err: 0.379, Validation loss: 0.6519261674880982  
Epoch 2: Train err: 0.366375, Train loss: 0.6400574240684509 | Validation err: 0.3595, Validation loss: 0.6199915554523469  
Epoch 3: Train err: 0.34225, Train loss: 0.6116521078348159 | Validation err: 0.3425, Validation loss: 0.636887101650238  
Epoch 4: Train err: 0.313375, Train loss: 0.5852544190883636 | Validation err: 0.3305, Validation loss: 0.597065006017685  
Epoch 5: Train err: 0.302875, Train loss: 0.5665943766832352 | Validation err: 0.3035, Validation loss: 0.5730490152835845  
Epoch 6: Train err: 0.279375, Train loss: 0.5459038650989533 | Validation err: 0.3055, Validation loss: 0.5840759282112121  
Epoch 7: Train err: 0.271875, Train loss: 0.5282641857862472 | Validation err: 0.308, Validation loss: 0.590907118320465  
Epoch 8: Train err: 0.259875, Train loss: 0.5090295715034008 | Validation err: 0.308, Validation loss: 0.5888441410064698  
Epoch 9: Train err: 0.250875, Train loss: 0.4967687402367592 | Validation err: 0.289, Validation loss: 0.5735659339427948  
Epoch 10: Train err: 0.23625, Train loss: 0.4753278249800205 | Validation err: 0.3005, Validation loss: 0.5866742620468139  
Epoch 11: Train err: 0.221125, Train loss: 0.4599813489615917 | Validation err: 0.2955, Validation loss: 0.6126594383716584  
Epoch 12: Train err: 0.21075, Train loss: 0.44328313660621643 | Validation err: 0.308, Validation loss: 0.6164555573463439  
Epoch 13: Train err: 0.202375, Train loss: 0.42531322610378264 | Validation err: 0.2925, Validation loss: 0.6418291474580765  
Epoch 14: Train err: 0.1875, Train loss: 0.4029668617844582 | Validation err: 0.307, Validation loss: 0.658977352142334  
Epoch 15: Train err: 0.174, Train loss: 0.38149635007977484 | Validation err: 0.3065, Validation loss: 0.6987648135423661  
Epoch 16: Train err: 0.1625, Train loss: 0.3601392722427845 | Validation err: 0.3045, Validation loss: 0.7223338420391083  
Epoch 17: Train err: 0.154, Train loss: 0.3441975792199373 | Validation err: 0.3275, Validation loss: 0.7796074187755585  
Epoch 18: Train err: 0.150625, Train loss: 0.3297359578534961 | Validation err: 0.322, Validation loss: 0.7965750460624694  
Epoch 19: Train err: 0.134875, Train loss: 0.3068180343657732 | Validation err: 0.3275, Validation loss: 0.8325750711560249  
Epoch 20: Train err: 0.12625, Train loss: 0.2907130775600672 | Validation err: 0.3365, Validation loss: 0.8838953652381897  
Epoch 21: Train err: 0.11225, Train loss: 0.271860674187541 | Validation err: 0.3295, Validation loss: 0.9316603469848633  
Epoch 22: Train err: 0.107875, Train loss: 0.2554583952873945 | Validation err: 0.325, Validation loss: 1.0496476212739945  
Epoch 23: Train err: 0.100625, Train loss: 0.24102731531858443 | Validation err: 0.334, Validation loss: 1.041201030910015  
Epoch 24: Train err: 0.099125, Train loss: 0.2457683519665152 | Validation err: 0.354, Validation loss: 1.0744450941085815  
Epoch 25: Train err: 0.0985, Train loss: 0.23269777764752508 | Validation err: 0.324, Validation loss: 1.1062232171297073  
Epoch 26: Train err: 0.0805, Train loss: 0.20908487829193473 | Validation err: 0.327, Validation loss: 1.2744335865974425  
Epoch 27: Train err: 0.089625, Train loss: 0.22396712187305093 | Validation err: 0.3615, Validation loss: 1.2231587070226668  
Epoch 28: Train err: 0.085625, Train loss: 0.21216214792104438 | Validation err: 0.3335, Validation loss: 1.3200735096931457  
Epoch 29: Train err: 0.07325, Train loss: 0.182615338915959 | Validation err: 0.3385, Validation loss: 1.48476123213768  
Epoch 30: Train err: 0.074625, Train loss: 0.18895039067883043 | Validation err: 0.3385, Validation loss: 1.48476123213768

dation err: 0.3395, Validation loss: 1.3933519666194916  
 Finished Training  
 Total time elapsed: 162.99 seconds

-----output curve for lar  
ge\_net



Answer:

With `batch_size = 16`, it takes 162.99 seconds. With `batch_size = 64`, it takes 106.67 seconds. With lower `batch_size`, the network takes more time to complete modeling.

As `batch_size` is representing the number of samples per batch. So, if by decreasing the `batch_size`, which means that the number of samples per batch will be decreased and the number of batches itself increases. So, it will take more time for each iteration and leads to an overall increasing of time.

We noticed by analysing the graph, by decreasing the `batch_size`, the training and validation value have larger separation for the trends, which represents it cannot avoid overfitting and enlarged the overfitting issue compared to the previous diagram.

Also, for the previous error/loss diagram, when epoch closer to 29 and `batch_size = 64`, the error/loss value are nearly 0.15 and 0.3. However, in this case, when epoch closer to 29 and `batch_size = 512`, the error/loss value are significantly higher than 0.3 and 1.4. Especially for the validation datasets loss, it has a significant loss, although the training sets has a lower error/loss value. Smaller amount of samples that has been calculated each iteration for the gradient and direction, which leads to the incorrectness for to precise to each training sample sets. According to the lecture notes, we optimize a different function loss at each iteration which leads to the trends to be too noisy and leads to the overfitting to the training sets and errorness for the validation sets.

## Part 4. Hyperparameter Search [6 pt]

### Part (a) - 2pt

Based on the plots from above, choose another set of values for the hyperparameters (`network`, `batch_size`, `learning_rate`) that you think would help you improve the validation accuracy. Justify your choice.

Answer:

From the comparison of the large network and small network from part two, we find out, the the large network have a better error/loss value and initially does not have the overfitting issue. So, it is more manageable and better network to be chosen.

From the adjustment on `batch_size` from part3, we noticed that, generally, we should avoid to use too small `batch_size`, since the gradients that calculated are based on small portion of the samples, which is not generalized for the model and leads to the overfitting issue. By increasing the `batch_size`, the overfitting issue will be solved. However, we should not increase it to a number that is too large, which will increase the overall loss/error values.

In the previous trials, the original trial is `batch_size = 64`, and we increased to `batch_size = 512`, which is too large. So, we can try to increase the `batch_size` to `batch_size = 128`, which is not that large number and try to see the result.

In conclusion, the set of values for the hyperparameters(`network`, `batch_size`, `learning_rate`) that I think would help me improve the validation accuracy is `network = large_net`, `batch_size = 128`, `learning_rate = 0.01`.

**Part (b) - 1pt**

Train the model with the hyperparameters you chose in part(a), and include the training curve.



In [ ]:

```
large_net = LargeNet()
train_net(large_net, 128, 0.01, 30) #Now, set batch_size=128

print("\n-----output curve for large_
net")
large_model_path = get_model_name("large", batch_size=128, learning_rate=0.01, e
poch=29)
plot_training_curve(large_model_path)
```

Files already downloaded and verified

Files already downloaded and verified



Epoch 1: Train err: 0.454375, Train loss: 0.6920222299439567 | Validation err: 0.426, Validation loss: 0.6897575631737709  
Epoch 2: Train err: 0.44475, Train loss: 0.6877915528085496 | Validation err: 0.411, Validation loss: 0.6820644438266754  
Epoch 3: Train err: 0.422125, Train loss: 0.681226979172419 | Validation err: 0.413, Validation loss: 0.6738783977925777  
Epoch 4: Train err: 0.410375, Train loss: 0.6737948031652541 | Validation err: 0.408, Validation loss: 0.6660139746963978  
Epoch 5: Train err: 0.39725, Train loss: 0.6661599912340679 | Validation err: 0.401, Validation loss: 0.6590253487229347  
Epoch 6: Train err: 0.376875, Train loss: 0.6567276139107961 | Validation err: 0.389, Validation loss: 0.6513239853084087  
Epoch 7: Train err: 0.374625, Train loss: 0.6488056731602502 | Validation err: 0.376, Validation loss: 0.6396529972553253  
Epoch 8: Train err: 0.359875, Train loss: 0.636491524794745 | Validation err: 0.37, Validation loss: 0.6430468559265137  
Epoch 9: Train err: 0.350625, Train loss: 0.628224181750464 | Validation err: 0.3475, Validation loss: 0.6243616007268429  
Epoch 10: Train err: 0.350125, Train loss: 0.6217665170866346 | Validation err: 0.3455, Validation loss: 0.6212388165295124  
Epoch 11: Train err: 0.33275, Train loss: 0.6107507698119633 | Validation err: 0.339, Validation loss: 0.6172070316970348  
Epoch 12: Train err: 0.32675, Train loss: 0.5993005765808953 | Validation err: 0.328, Validation loss: 0.60542818531394  
Epoch 13: Train err: 0.3215, Train loss: 0.5949317689925905 | Validation err: 0.319, Validation loss: 0.6001784466207027  
Epoch 14: Train err: 0.30875, Train loss: 0.5798390091411652 | Validation err: 0.3295, Validation loss: 0.6141054034233093  
Epoch 15: Train err: 0.3055, Train loss: 0.5710504211130596 | Validation err: 0.3365, Validation loss: 0.6149555072188377  
Epoch 16: Train err: 0.296, Train loss: 0.5657848819853768 | Validation err: 0.3335, Validation loss: 0.6067412309348583  
Epoch 17: Train err: 0.29375, Train loss: 0.5584587341263181 | Validation err: 0.313, Validation loss: 0.597804632037878  
Epoch 18: Train err: 0.28325, Train loss: 0.549594487462725 | Validation err: 0.306, Validation loss: 0.5869672819972038  
Epoch 19: Train err: 0.27475, Train loss: 0.5373550625074477 | Validation err: 0.323, Validation loss: 0.5953679867088795  
Epoch 20: Train err: 0.274375, Train loss: 0.5338996487950521 | Validation err: 0.302, Validation loss: 0.5937820374965668  
Epoch 21: Train err: 0.266, Train loss: 0.5258608611803206 | Validation err: 0.302, Validation loss: 0.5821750834584236  
Epoch 22: Train err: 0.261875, Train loss: 0.5189807481235928 | Validation err: 0.309, Validation loss: 0.5881347022950649  
Epoch 23: Train err: 0.261125, Train loss: 0.5148856072198778 | Validation err: 0.3025, Validation loss: 0.5825694780796766  
Epoch 24: Train err: 0.256, Train loss: 0.5103520902376326 | Validation err: 0.317, Validation loss: 0.6025396026670933  
Epoch 25: Train err: 0.24875, Train loss: 0.50384506961656 | Validation err: 0.3005, Validation loss: 0.5932902377098799  
Epoch 26: Train err: 0.2435, Train loss: 0.4893661959776803 | Validation err: 0.3045, Validation loss: 0.5805531851947308  
Epoch 27: Train err: 0.241125, Train loss: 0.4899706646563515 | Validation err: 0.3105, Validation loss: 0.600347001105547  
Epoch 28: Train err: 0.239875, Train loss: 0.4822759576260097 | Validation err: 0.3045, Validation loss: 0.5907586440443993  
Epoch 29: Train err: 0.230375, Train loss: 0.47525750786539106 | Validation err: 0.32, Validation loss: 0.6141728721559048  
Epoch 30: Train err: 0.22825, Train loss: 0.4687656448000953 | Validation

tion err: 0.3095, Validation loss: 0.5970310214906931

Finished Training

Total time elapsed: 119.29 seconds

-----output curve for lar  
ge\_net



## Part (c) - 2pt

Based on your result from Part(a), suggest another set of hyperparameter values to try. Justify your choice.

Answer:

Additional to the analysis in part a), we can also make some adjustment towards the `learning_rate` and `epoch_size`.

From the adjustment on `learning_rate` from part3, we noticed that, generally, we should avoid the learning rate to be very large, since the step size when updating the weights each time is too large, which is not generalized for the model and leads to the overfitting issue. By decreasing the `learning_rate`, the overfitting issue will be solved. However, we should not decreased it to a number that is too small, which will increase the overall loss/error values. In the previous trials, the original trial is `learning_rate = 0.01` and we decreqased it to `learning_rate = 0.001`, which is too small. So, we can try to decrease the `learning_rate` to `learning_rate = 0.005`, which is not that small number and try to see the result.

In conclusion, the set of values for the hyperparameters(`network`, `batch_size`, `learning_rate`) that I think would help me improve the validation accuracy is `network = large_net`, `batch_size = 128`, `learning_rate = 0.005`.

## Part (d) - 1pt

Train the model with the hyperparameters you chose in part(c), and include the training curve.

In [ ]:

```
large_net = LargeNet()
train_net(large_net, 128, 0.005, 30) #Now, set batch_size=128, learning_rate = 0.005

print("\n-----output curve for large_net")
large_model_path = get_model_name("large", batch_size=128, learning_rate=0.005, epoch=29)
plot_training_curve(large_model_path)
```

Files already downloaded and verified

Files already downloaded and verified



Epoch 1: Train err: 0.466625, Train loss: 0.6925613236805749 | Validation err: 0.4305, Validation loss: 0.6916250362992287  
Epoch 2: Train err: 0.45025, Train loss: 0.6910346017943488 | Validation err: 0.4295, Validation loss: 0.6889704614877701  
Epoch 3: Train err: 0.4305, Train loss: 0.6885886050405956 | Validation err: 0.4165, Validation loss: 0.6850010603666306  
Epoch 4: Train err: 0.430625, Train loss: 0.6850024423901997 | Validation err: 0.4125, Validation loss: 0.6797147057950497  
Epoch 5: Train err: 0.420875, Train loss: 0.6813881851377941 | Validation err: 0.411, Validation loss: 0.6762721762061119  
Epoch 6: Train err: 0.415875, Train loss: 0.6773001438095456 | Validation err: 0.4125, Validation loss: 0.672969501465559  
Epoch 7: Train err: 0.405, Train loss: 0.6732292080682422 | Validation err: 0.4055, Validation loss: 0.6696001030504704  
Epoch 8: Train err: 0.4, Train loss: 0.6694153556748043 | Validation err: 0.402, Validation loss: 0.6651207581162453  
Epoch 9: Train err: 0.3905, Train loss: 0.6657110272891937 | Validation err: 0.3965, Validation loss: 0.6595559529960155  
Epoch 10: Train err: 0.38475, Train loss: 0.6609666309659443 | Validation err: 0.398, Validation loss: 0.6561728976666927  
Epoch 11: Train err: 0.374625, Train loss: 0.6547484369505019 | Validation err: 0.3895, Validation loss: 0.6526741832494736  
Epoch 12: Train err: 0.379125, Train loss: 0.6503025397421822 | Validation err: 0.382, Validation loss: 0.644683513790369  
Epoch 13: Train err: 0.3685, Train loss: 0.6430244833704025 | Validation err: 0.377, Validation loss: 0.6408736705780029  
Epoch 14: Train err: 0.35375, Train loss: 0.634985539648268 | Validation err: 0.3715, Validation loss: 0.6393358670175076  
Epoch 15: Train err: 0.35375, Train loss: 0.6319066002255395 | Validation err: 0.385, Validation loss: 0.6489016860723495  
Epoch 16: Train err: 0.35125, Train loss: 0.6263072131172059 | Validation err: 0.357, Validation loss: 0.6286376938223839  
Epoch 17: Train err: 0.349875, Train loss: 0.6246873547160436 | Validation err: 0.351, Validation loss: 0.6268594488501549  
Epoch 18: Train err: 0.343, Train loss: 0.6182309228276449 | Validation err: 0.3415, Validation loss: 0.6208269745111465  
Epoch 19: Train err: 0.33675, Train loss: 0.6112786872046334 | Validation err: 0.348, Validation loss: 0.6215575821697712  
Epoch 20: Train err: 0.334875, Train loss: 0.6074399777821132 | Validation err: 0.3385, Validation loss: 0.6121908687055111  
Epoch 21: Train err: 0.328375, Train loss: 0.603059540665339 | Validation err: 0.337, Validation loss: 0.6123869866132736  
Epoch 22: Train err: 0.327375, Train loss: 0.5975976425503927 | Validation err: 0.333, Validation loss: 0.6109451837837696  
Epoch 23: Train err: 0.322375, Train loss: 0.5930433736907111 | Validation err: 0.3245, Validation loss: 0.6026575490832329  
Epoch 24: Train err: 0.317875, Train loss: 0.5886693076481895 | Validation err: 0.327, Validation loss: 0.5991024523973465  
Epoch 25: Train err: 0.306375, Train loss: 0.5787286261717478 | Validation err: 0.315, Validation loss: 0.596717856824398  
Epoch 26: Train err: 0.30725, Train loss: 0.572699801316337 | Validation err: 0.321, Validation loss: 0.5967751257121563  
Epoch 27: Train err: 0.302875, Train loss: 0.5695816690013522 | Validation err: 0.3185, Validation loss: 0.5978134162724018  
Epoch 28: Train err: 0.300875, Train loss: 0.5655642464047387 | Validation err: 0.334, Validation loss: 0.6142176277935505  
Epoch 29: Train err: 0.294125, Train loss: 0.5602825222507356 | Validation err: 0.3275, Validation loss: 0.6087971664965153  
Epoch 30: Train err: 0.291875, Train loss: 0.5594303224767957 | Validation err: 0.3275, Validation loss: 0.6087971664965153

ation err: 0.328, Validation loss: 0.6033484600484371

Finished Training

Total time elapsed: 116.50 seconds

-----output curve for lar  
ge\_net



## Part 5. Evaluating the Best Model [15 pt]

### Part (a) - 1pt

Choose the **best** model that you have so far. This means choosing the best model checkpoint, including the choice of `small_net` vs `large_net`, the `batch_size`, `learning_rate`, **and the epoch number**.

Modify the code below to load your chosen set of weights to the model object `net`.

Answer:

Based on the graph in part 4 b), as epoch\_number = 26, the diagram reached a quite low error/loss value. So, we decreased the epoch\_number to 26 to avoid the large variation from 26 to 30.

In conclusion, the best model that I can provide so far is: `network = large_net` , `batch_size = 128` , `learning_rate = 0.01` and `epoch_number = 26` .



In [ ]:

```
net = LargeNet()
train_net(net, 128, 0.01, 26)
model_path = get_model_name(net.name, batch_size=128, learning_rate=0.01, epoch=
25)
state = torch.load(model_path)
net.load_state_dict(state)
```

```
Files already downloaded and verified
Files already downloaded and verified
Epoch 1: Train err: 0.454375, Train loss: 0.6920222299439567 |Validation err: 0.426, Validation loss: 0.6897575631737709
Epoch 2: Train err: 0.44475, Train loss: 0.6877915528085496 |Validation err: 0.411, Validation loss: 0.6820644438266754
Epoch 3: Train err: 0.422125, Train loss: 0.681226979172419 |Validation err: 0.413, Validation loss: 0.6738783977925777
Epoch 4: Train err: 0.410375, Train loss: 0.6737948031652541 |Validation err: 0.408, Validation loss: 0.6660139746963978
Epoch 5: Train err: 0.39725, Train loss: 0.6661599912340679 |Validation err: 0.401, Validation loss: 0.6590253487229347
Epoch 6: Train err: 0.376875, Train loss: 0.6567276139107961 |Validation err: 0.389, Validation loss: 0.6513239853084087
Epoch 7: Train err: 0.374625, Train loss: 0.6488056731602502 |Validation err: 0.376, Validation loss: 0.6396529972553253
Epoch 8: Train err: 0.359875, Train loss: 0.636491524794745 |Validation err: 0.37, Validation loss: 0.6430468559265137
Epoch 9: Train err: 0.350625, Train loss: 0.628224181750464 |Validation err: 0.3475, Validation loss: 0.6243616007268429
Epoch 10: Train err: 0.350125, Train loss: 0.6217665170866346 |Validation err: 0.3455, Validation loss: 0.6212388165295124
Epoch 11: Train err: 0.33275, Train loss: 0.6107507698119633 |Validation err: 0.339, Validation loss: 0.6172070316970348
Epoch 12: Train err: 0.32675, Train loss: 0.5993005765808953 |Validation err: 0.328, Validation loss: 0.60542818531394
Epoch 13: Train err: 0.3215, Train loss: 0.5949317689925905 |Validation err: 0.319, Validation loss: 0.6001784466207027
Epoch 14: Train err: 0.30875, Train loss: 0.5798390091411652 |Validation err: 0.3295, Validation loss: 0.6141054034233093
Epoch 15: Train err: 0.3055, Train loss: 0.5710504211130596 |Validation err: 0.3365, Validation loss: 0.6149555072188377
Epoch 16: Train err: 0.296, Train loss: 0.5657848819853768 |Validation err: 0.3335, Validation loss: 0.6067412309348583
Epoch 17: Train err: 0.29375, Train loss: 0.5584587341263181 |Validation err: 0.313, Validation loss: 0.597804632037878
Epoch 18: Train err: 0.28325, Train loss: 0.549594487462725 |Validation err: 0.306, Validation loss: 0.5869672819972038
Epoch 19: Train err: 0.27475, Train loss: 0.5373550625074477 |Validation err: 0.323, Validation loss: 0.5953679867088795
Epoch 20: Train err: 0.274375, Train loss: 0.5338996487950521 |Validation err: 0.302, Validation loss: 0.5937820374965668
Epoch 21: Train err: 0.266, Train loss: 0.5258608611803206 |Validation err: 0.302, Validation loss: 0.5821750834584236
Epoch 22: Train err: 0.261875, Train loss: 0.5189807481235928 |Validation err: 0.309, Validation loss: 0.5881347022950649
Epoch 23: Train err: 0.261125, Train loss: 0.5148856072198778 |Validation err: 0.3025, Validation loss: 0.5825694780796766
Epoch 24: Train err: 0.256, Train loss: 0.5103520902376326 |Validation err: 0.317, Validation loss: 0.6025396026670933
Epoch 25: Train err: 0.24875, Train loss: 0.50384506961656 |Validation err: 0.3005, Validation loss: 0.5932902377098799
Epoch 26: Train err: 0.2435, Train loss: 0.4893661959776803 |Validation err: 0.3045, Validation loss: 0.5805531851947308
Finished Training
Total time elapsed: 100.65 seconds
```

Out[ ]:


<All keys matched successfully>

## Part (b) - 2pt

Justify your choice of model from part (a).

In [ ]:

```
plot_training_curve(model_path)
```



Answer:

As we can see from the final output for loss/error values are for our model `net_name=large_net`, `batch_size=128`, `learning_rate=0.01`, `epoch=26` is Validation err: 0.3045, Validation loss: 0.5805531851947308 which is lower in err/loss than the default training values `net_name=large_net`, `batch_size=64`, `learning_rate=0.01`, `epoch=30`, which is Validation err: 0.317, Validation loss: 0.7528756745159626 .

Both the error/loss value and the trace of the diagram from smaller epoch to larger epoch is better than the default one. Since, by increasing the `batch_size` and decreasing the `learning_rate`, the overfitting of the model has been improved. So, the model for part a will have an overall better performance than the default one, as the trend of validation error is slightly higher than the training error.

## Part (c) - 2pt

Using the code in Part 0, any code from lecture notes, or any code that you write, compute and report the **test classification error** for your chosen model.

In [ ]:

```
# If you use the `evaluate` function provided in part 0, you will need to
# set batch_size > 1
train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes=["cat", "dog"],
    batch_size=128)

criterion = nn.BCEWithLogitsLoss()
train_err, train_loss = evaluate(net, train_loader, criterion)
val_err, val_loss = evaluate(net, val_loader, criterion)
test_err, test_loss = evaluate(net, test_loader, criterion)

print("training: err = " + str(train_err) + " / loss = " + str(train_loss))
print("validation: err = " + str(val_err) + " / loss = " + str(val_loss))
print("testing: err = " + str(test_err) + " / loss = " + str(test_loss))
```

```
Files already downloaded and verified
Files already downloaded and verified
training: err = 0.2345 / loss = 0.47904046518462046
validation: err = 0.3045 / loss = 0.5845830701291561
testing: err = 0.29 / loss = 0.5610999148339033
```

## Part (d) - 3pt

How does the test classification error compare with the **validation error**? Explain why you would expect the test error to be *higher* than the validation error.

Answer:

As we can see from the output in part c, the error value for testing is smaller than the error value for validation.

---

However, we should expect the test error to be higher than the validation error. Because the reason that there is a set of data for validation is to adjust the hyperparameters, to improve the model to reach a good fit, as we use the training data to build up and train the model, and testing data to test the result. So, as we already used the validation set while adjusting the model towards a good way, so the error will be expected to be less than a group of data that has not been seen for testing. But there exist some exceptions is that, if the data sets is not randomized enough and the testing images is more similar to the training sets, the error of the testing dataset will be lower than fully randomized situation.

## Part (e) - 2pt

Why did we only use the test data set at the very end? Why is it important that we use the test data as little as possible?

The reason why we should use the test data set at the very end for testing the model for overall performance is that, as the testing data sets are never used in training or adjusting hyperparameters. So, compared to other datasets that contribute to the overall building of the model, the testing dataset is unbiased and is good for the overall testing for the performance at the very end, when the model is completed.

---

The reason why it is important that we use the test data as little as possible is that, as we evaluate the model each time, the data input and labels will affect the model. So, if we used the test data too many times, the testing data will be less unbiased, and even become biased in the end. If the testing dataset becomes biased, then the error/loss after the evaluation is not valuable for us to see the overall performance of the model.

## Part (f) - 5pt

How does your best CNN model compare with a 2-layer ANN model (no convolutional layers) on classifying cat and dog images. You can use a 2-layer ANN architecture similar to what you used in Lab 1. You should explore different hyperparameter settings to determine how well you can do on the validation dataset. Once satisfied with the performance, you may test it out on the test data.

Hint: The ANN in lab 1 was applied on grayscale images. The cat and dog images are colour (RGB) and so you will need to flatten and concatenate all three colour layers before feeding them into an ANN.

In [ ]:

```
# copied and modified the code from lab 1 to classifying cat and dog images

import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import datasets, transforms
import matplotlib.pyplot as plt # for plotting
import torch.optim as optim

torch.manual_seed(1) # set the random seed

# define a 2-layer artificial neural network
class Pigeon(nn.Module):
    def __init__(self):
        super(Pigeon, self).__init__()
        self.layer1 = nn.Linear(3*32*32, 30) #image size 32*32, RGB is 3
        self.layer2 = nn.Linear(30, 1)
        self.name = "pigeon for ANN"
    def forward(self, img):
        flattened = img.view(-1, 3*32*32)
        activation1 = self.layer1(flattened)
        activation1 = F.relu(activation1)
        activation2 = self.layer2(activation1)
        return activation2.squeeze(1)

pigeon = Pigeon()
train_net(pigeon, 128, 0.01, 26)
ANN_model_path = get_model_name("pigeon for ANN", batch_size=128, learning_rate=
0.01, epoch=25)
plot_training_curve(ANN_model_path)

pigeon1 = Pigeon()
train_net(pigeon1, 64, 0.01, 30)
ANN_model_path1 = get_model_name("pigeon for ANN", batch_size=64, learning_rate=
0.01, epoch=29)
plot_training_curve(ANN_model_path1)
```

Files already downloaded and verified

Files already downloaded and verified

Epoch 1: Train err: 0.41775, Train loss: 0.668205031326839 |Validation err: 0.399, Validation loss: 0.6558435969054699

Epoch 2: Train err: 0.371125, Train loss: 0.6407988799942864 |Validation err: 0.39, Validation loss: 0.6511164344847202

Epoch 3: Train err: 0.355375, Train loss: 0.628787152350895 |Validation err: 0.383, Validation loss: 0.645381610840559

Epoch 4: Train err: 0.345, Train loss: 0.6168104429093618 |Validation err: 0.399, Validation loss: 0.6566037982702255

Epoch 5: Train err: 0.333125, Train loss: 0.6075799654400538 |Validation err: 0.392, Validation loss: 0.6480461545288563

Epoch 6: Train err: 0.32125, Train loss: 0.597667814247192 |Validation err: 0.383, Validation loss: 0.6525003537535667

Epoch 7: Train err: 0.31075, Train loss: 0.5869401277057709 |Validation err: 0.3715, Validation loss: 0.649830024689436

Epoch 8: Train err: 0.29775, Train loss: 0.5727507538265653 |Validation err: 0.372, Validation loss: 0.6525772772729397

Epoch 9: Train err: 0.292875, Train loss: 0.5635531393308488 |Validation err: 0.381, Validation loss: 0.6523666307330132

Epoch 10: Train err: 0.285, Train loss: 0.5517164189664144 |Validation err: 0.367, Validation loss: 0.6559775620698929

Epoch 11: Train err: 0.27525, Train loss: 0.5387107602187565 |Validation err: 0.371, Validation loss: 0.6658255383372307

Epoch 12: Train err: 0.26125, Train loss: 0.5243138268826499 |Validation err: 0.3685, Validation loss: 0.6797713413834572

Epoch 13: Train err: 0.265875, Train loss: 0.5235499328091031 |Validation err: 0.388, Validation loss: 0.7016571462154388

Epoch 14: Train err: 0.25325, Train loss: 0.5102075560698434 |Validation err: 0.39, Validation loss: 0.6870112046599388

Epoch 15: Train err: 0.2395, Train loss: 0.49323967524937223 |Validation err: 0.364, Validation loss: 0.6964708790183067

Epoch 16: Train err: 0.23175, Train loss: 0.4856313954270075 |Validation err: 0.3535, Validation loss: 0.691399596631527

Epoch 17: Train err: 0.229375, Train loss: 0.47713576943155317 |Validation err: 0.3605, Validation loss: 0.6826366521418095

Epoch 18: Train err: 0.216875, Train loss: 0.45949133710255696 |Validation err: 0.377, Validation loss: 0.7249729223549366

Epoch 19: Train err: 0.194875, Train loss: 0.4395922715701754 |Validation err: 0.378, Validation loss: 0.7140126973390579

Epoch 20: Train err: 0.197, Train loss: 0.43045873821727815 |Validation err: 0.368, Validation loss: 0.7396925576031208

Epoch 21: Train err: 0.1995, Train loss: 0.4256557420132652 |Validation err: 0.3735, Validation loss: 0.7555751167237759

Epoch 22: Train err: 0.18025, Train loss: 0.40419161319732666 |Validation err: 0.3675, Validation loss: 0.7680577374994755

Epoch 23: Train err: 0.175625, Train loss: 0.394673364503043 |Validation err: 0.364, Validation loss: 0.7445585243403912



Epoch 24: Train err: 0.17275, Train loss: 0.3944419977210817 |Validation err: 0.378, Validation loss: 0.7795960195362568

Epoch 25: Train err: 0.17225, Train loss: 0.38616057284294614 |Validation err: 0.3625, Validation loss: 0.77102230489254

Epoch 26: Train err: 0.149, Train loss: 0.35587310459878707 |Validation err: 0.36, Validation loss: 0.8014938458800316

Finished Training

Total time elapsed: 73.28 seconds





Files already downloaded and verified

Files already downloaded and verified

Epoch 1: Train err: 0.40675, Train loss: 0.6628824081420899 | Validation err: 0.383, Validation loss: 0.6505479384213686  
Epoch 2: Train err: 0.374, Train loss: 0.6387409133911133 | Validation err: 0.394, Validation loss: 0.6582276318222284  
Epoch 3: Train err: 0.3595, Train loss: 0.6286095876693726 | Validation err: 0.379, Validation loss: 0.6436176169663668  
Epoch 4: Train err: 0.34775, Train loss: 0.6167842669486999 | Validation err: 0.4055, Validation loss: 0.6707729957997799  
Epoch 5: Train err: 0.333125, Train loss: 0.6061416540145874 | Validation err: 0.3765, Validation loss: 0.6455826908349991  
Epoch 6: Train err: 0.314, Train loss: 0.5921389698982239 | Validation err: 0.3795, Validation loss: 0.6513363718986511  
Epoch 7: Train err: 0.30875, Train loss: 0.5826152532100678 | Validation err: 0.379, Validation loss: 0.6588644441217184  
Epoch 8: Train err: 0.289875, Train loss: 0.5654429728984833 | Validation err: 0.3615, Validation loss: 0.6509492415934801  
Epoch 9: Train err: 0.292, Train loss: 0.5568346982002258 | Validation err: 0.384, Validation loss: 0.6689472999423742  
Epoch 10: Train err: 0.27475, Train loss: 0.5397703149318696 | Validation err: 0.3755, Validation loss: 0.6600630935281515  
Epoch 11: Train err: 0.269, Train loss: 0.5283022401332855 | Validation err: 0.371, Validation loss: 0.6803382877260447  
Epoch 12: Train err: 0.251375, Train loss: 0.5090102548599243 | Validation err: 0.38, Validation loss: 0.7129230536520481  
Epoch 13: Train err: 0.239625, Train loss: 0.49473757433891297 | Validation err: 0.3635, Validation loss: 0.7111866381019354  
Epoch 14: Train err: 0.230375, Train loss: 0.4838444561958313 | Validation err: 0.378, Validation loss: 0.7431046180427074  
Epoch 15: Train err: 0.21875, Train loss: 0.4653008215427399 | Validation err: 0.3715, Validation loss: 0.7926676496863365  
Epoch 16: Train err: 0.21325, Train loss: 0.45780208277702333 | Validation err: 0.367, Validation loss: 0.7463004644960165  
Epoch 17: Train err: 0.20325, Train loss: 0.4411180346012116 | Validation err: 0.374, Validation loss: 0.7232418414205313  
Epoch 18: Train err: 0.1895, Train loss: 0.4146708402633667 | Validation err: 0.3735, Validation loss: 0.772447507828474  
Epoch 19: Train err: 0.190125, Train loss: 0.4179953124523163 | Validation err: 0.389, Validation loss: 0.9067568760365248  
Epoch 20: Train err: 0.188125, Train loss: 0.4157635028362274 | Validation err: 0.3625, Validation loss: 0.7911309562623501  
Epoch 21: Train err: 0.16475, Train loss: 0.3724208711385727 | Validation err: 0.3745, Validation loss: 0.8082200028002262  
Epoch 22: Train err: 0.167125, Train loss: 0.3893840044736862 | Validation err: 0.38, Validation loss: 0.9257937390357256  
Epoch 23: Train err: 0.15225, Train loss: 0.34655749118328094 | Validation err: 0.36, Validation loss: 0.834901224821806  
Epoch 24: Train err: 0.15575, Train loss: 0.36253183376789094 | Validation err: 0.3765, Validation loss: 0.8323851712048054  
Epoch 25: Train err: 0.144375, Train loss: 0.33305981373786925 | Validation err: 0.364, Validation loss: 0.8360881507396698  
Epoch 26: Train err: 0.1295, Train loss: 0.31471012794971465 | Validation err: 0.3755, Validation loss: 0.9483208116143942  
Epoch 27: Train err: 0.1375, Train loss: 0.32656157648563383 | Validation err: 0.3605, Validation loss: 0.9301714859902859  
Epoch 28: Train err: 0.129625, Train loss: 0.3052001684904099 | Validation err: 0.365, Validation loss: 0.9871594347059727  
Epoch 29: Train err: 0.109, Train loss: 0.27064790081977846 | Validation err: 0.3715, Validation loss: 1.0711834449321032  
Epoch 30: Train err: 0.12175, Train loss: 0.29704382932186124 | Validation

ation err: 0.3685, Validation loss: 1.0325132552534342  
Finished Training  
Total time elapsed: 94.88 seconds



In [ ]:

```

train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes=["cat", "dog"],
    batch_size=128)

cirterion = nn.BCEWithLogitsLoss()
ANN_train_err, ANN_train_loss = evaluate(pigeon, train_loader, cirterion)
ANN_val_err, ANN_val_loss = evaluate(pigeon, val_loader, cirterion)
ANN_test_err, ANN_test_loss = evaluate(pigeon, test_loader, cirterion)

print("ANN training: err = " + str(ANN_train_err) + " / loss = " + str(ANN_train_
loss))
print("ANN validation: err = " + str(ANN_val_err) + " / loss = " + str(ANN_val_lo
ss))
print("ANN testing: err = " + str(ANN_test_err) + " / loss = " + str(ANN_test_lo
s))

print("\n")
train_loader1, val_loader1, test_loader1, classes1 = get_data_loader(
    target_classes=["cat", "dog"],
    batch_size=64)

cirterion1 = nn.BCEWithLogitsLoss()
ANN_train_err1, ANN_train_loss1 = evaluate(pigeon1, train_loader1, cirterion1)
ANN_val_err1, ANN_val_loss1 = evaluate(pigeon1, val_loader1, cirterion1)
ANN_test_err1, ANN_test_loss1 = evaluate(pigeon1, test_loader1, cirterion1)

print("ANN1 training: err = " + str(ANN_train_err1) + " / loss = " + str(ANN_trai
n_loss1))
print("ANN1 validation: err = " + str(ANN_val_err1) + " / loss = " + str(ANN_val_
loss1))
print("ANN1 testing: err = " + str(ANN_test_err1) + " / loss = " + str(ANN_test_l
oss1))

```

Files already downloaded and verified

Files already downloaded and verified

ANN training: err = 0.131875 / loss = 0.3219138686619108

ANN validation: err = 0.36 / loss = 0.8034975901246071

ANN testing: err = 0.366 / loss = 0.8008576855063438

Files already downloaded and verified

Files already downloaded and verified

ANN1 training: err = 0.096875 / loss = 0.2417677720785141

ANN1 validation: err = 0.3685 / loss = 1.0345731191337109

ANN1 testing: err = 0.3775 / loss = 1.0677751675248146

Answer:

As we discovered, the first training set for ANN is with the hyperparameters that I choosed, and the second training set for ANN is with the hyperparameters as default. We can see from the graph, by using the chosen hyperparameters, it improves a little bit of the overfitting issue for ANN model. However, in general comparing the error and lost, ANN is not a good model for the performance, because it has larger error and loss than the CNN model.